

Srinivas Dhanwada, Collin McIntyre, Matt Wall, Ben Weno

Overview

- L-SAP
- Problems with the Existing Verifier Pipeline
- Our Solution
- Patch Algorithm
- Differencing Algorithm
- Website Redesign
- Automation Design
- Testing & Evaluation Plan
- Project Schedule & Management
- Summary

L-SAP

L-SAP

- Scalable and Accurate Lock/Unlock Pairing for the Linux Kernel
- Looks for Lock/Unlock pairing inconsistencies throughout the kernel codebase
- Built on top of Atlas plugin to provide model-based verification
 - Other verification tools are unable to provide this for locking/unlocking problems
 - Example:
 - Lock a variable
 - Set it as a global variable
 - Unlock the global variable later in the code path.
- L-SAP Solves this problem by translating the code path into models and graphs
 - For each locking instance:
 - A graph is created tracking all the possible methods it can reach
 - A graph is created for each method tracking the status of the locking instance
 - L-SAP provides a conclusion, but also provides these graphs as evidence
 - Humans can look at this evidence and analyze inconclusive cases faster than looking at the codebase directly

Problems with the Existing Verifier Pipeline

Existing Verifier Pipeline Problems

- Currently L-SAP is unable to keep up with the pace of Linux development:
 - Have to manually create patch for each version
 - Have to manually upload results to website
- No existing way to map instances between different versions
 - How can we make sure a lock mismatch in a previous version was actually fixed?
- Current website is difficult to use
 - Consists of two lists for two different types of locking that each have thousands of links
 - No search tools available
 - Difficult to find mismatched locks

Requirements For Our Solution

Functional:

- Automatically recognize kernel updates and start the tool
- Create and apply a patch for the tool for each new version of the kernel
- Run the tool with the new patch
- Create a difference mapping for each locking instance
- Post the results to the website
 - The website must properly display the new results in an organized manner

Non Functional

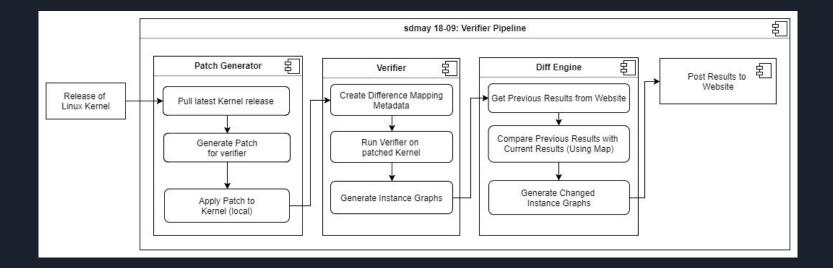
- The tool must run in a reasonable amount of time
- The website must have a high level of usability, and scale to it's demands

Our Solution

Our Solution

- Modularize and automate the pipeline
 - Patch Creation and Application
 - Instance Mapping
 - Running L-SAP
 - Differencing Results
- Redesign website to make tracking results easier
 - Provide search criteria
 - Text search
 - By driver search
 - Clearly show which instances are mismatched
 - Show the difference from previous versions to verify updates in the kernel quicker

Our Solution



Patch Algorithm

Patching Algorithm

• What is the patch?

 Redirection of all locking function calls to a single function per lock type

• Why is it needed?

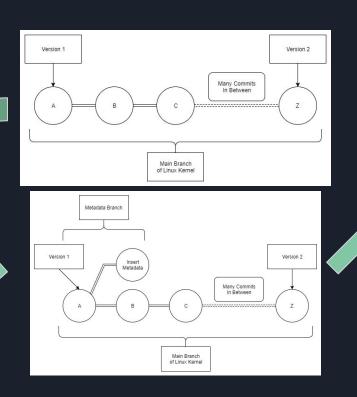
- L-SAP matches locks and unlocks based on function calls
- Redirecting locking calls to a single function reduces computation time
- Since kernel structure is all we care about, we can redirect to an empty function

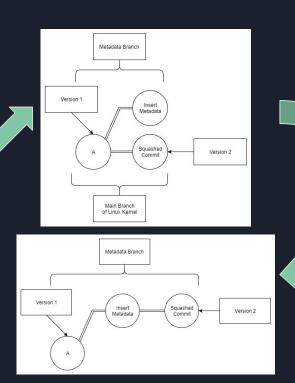


Patching Algorithm

- Before the algorithm runs
 - Define parameters for mutex lock functions/macros
 - Define parameters for spinlock functions/macros
 - Determine kernel location
- Overview
 - Get function/macro information from existing headers
 - Generate patched header files
 - Using function/macro information...
 - Remove existing function/macro declarations
 - Remove existing function/macro implementations
- Process will be the same for mutex locks and spinlocks

- Problem: Current Implementation of L-SAP assigns a random ID to each instance
 - No mapping between versions is possible, which means no comparison is possible
 - Looking at the same point in source code does not work either
 - Additions/Deletions could happen nearby causing a shift in the location of the instance
 - Changes to the locking variable name could also make it hard to find between two versions
 - Need some way to insert data within source code and have it move as the source code moves.
- Solution: Leverage the Linux Kernel's use of Git to help with tracking the differences between versions!
 - Insert a comment tag at the location of the instance -- commit these to a branch off of the first version
 - Use Git to rebase the differences between version onto the new branch
 - This inserts the comments first, then applies all the changes inserting the comments and having them move with the source code!





- Mapping between different versions has been created
- Need to fetch data and look for differences between results
 - This can be done by running the verifier on our modified version with metadata tags
 - Look at each instance for the metadata tag
 - Export these differences into a spreadsheet
- Due to the output of L-SAP, the only metric of differencing right now is the final status
 - o An update to L-SAP might expose more data about the graph structures we can use
 - We will be able to compare the number of edges/nodes in a graph as well as the connections to see
 if the graph has changed even with the same status

Website Redesign

Angular + Typescript

PROS CONS

- Modularity
- Data Binding
- Cleaner Implementation
- Optimization
- Maintenance

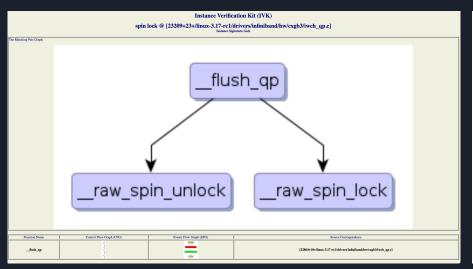
- Harder to implement than static content
- Some server side implementation required
- Must have javascript enabled
- One entry point to the site

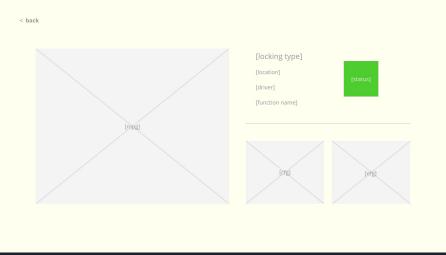
Home Page Design

| Spin & Mutex Verification Instances for Linux Kernel (3.17-rc1) | | | | | | | | | |
|---|--|--|--|--|--|--|--|--|--|
| Spin Locks | Mutex Locks | | | | | | | | |
| 1. [23209±23+/linux.3.17-re.l/drivers/infinil/band/hw/cxgb3/ivech_qp_c] 2. [25709±24+/linux.3.17-re.l/drivers/infinil/band/hw/cxgb3/ivech_qp_c] 3. [21399±41+/linux.3.17-re.l/drivers/stateging/siter/flue/klinds/c2lbind/c2lbind_cb_c] 3. [21399±41+/linux.3.17-re.l/drivers/stateging/siter/flue/klinds/c2lbind/c2lbind_cb_c] 4. [2121548+fl-il/minux.3.17-re.l/drivers/stategil/stategil/secale_cd_c] 5. [212958+fl-il/minux.3.17-re.l/drivers/stategil/secale_cd_c] 6. [231628-77-il/minux.3.17-re.l/drivers/stategil/secale_cd_c] 7. [610014-65-vlinux.3.17-re.l/drivers/stategil/secale_cd_c] 8. [16637-84-vlinux.3.17-re.l/drivers/stategil/secale_cd_c] 9. [234001+99-vlinux.3.17-re.l/drivers/stategil/secale_cd_c] 11. [881528-82-vlinux.3.17-re.l/drivers/stategil/secale_cd_c] 12. [74215-89-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 13. [7409-9-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 14. [2876-44-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 15. [8693-40-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 16. [8693-40-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 17. [8404-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 18. [8693-40-vlinux.3.17-re.l/drivers/stategil/secale_sac_cd_c] 18. [8693-40-vlinux.3.17-re.l/drivers/stategil/secale_cd_c] 18. [8693-40-vlinux.3.17-re.l/drivers/stategil/secale_cd_cd_cd_cd_cd_cd_cd_cd_cd_cd_cd_cd_cd_ | Mutex Locks | | | | | | | | |
| 28. [12461+26+/linux-3.17-rc1/drivers/staging/lustre/lustre/ptlrpc/nrs.c] | 28. [15135+32+/linux-3.17-rc1/drivers/tty/n_tty.c] | | | | | | | | |
| [12391+29+/linux-3.17-rc1/drivers/staging/lustre/lustre/mgc/mgc_request.c] [47025+49+/linux-3.17-rc1/drivers/scsi/mpt2sas/mpt2sas_scsih.c] | 29. [23777+27+/linux-3.17-rc1/drivers/tty/moxa.c] 30. [10637+31+/linux-3.17-rc1/drivers/hymnon/asc7621.c] | | | | | | | | |
| 31. [45586+22+/linux-3.17-rc1/drivers/usb/gadget/udc/dummy_hcd.c] | 31. [60076+32+/linux-3.17-rc1/drivers/scsi/3w-9xxx.c] | | | | | | | | |
| 32. [15173+42+/linux-3.17-rc1/drivers/misc/hpilo.c] | 32. [6025+26+/linux-3.17-rc1/drivers/pcmcia/pcmcia resource.c] | | | | | | | | |
| 33. [20999+54+/linux-3.17-rc1/drivers/staging/comedi/drivers/ni_660x.c] | 33. [10031+23+/linux-3.17-rc1/drivers/media/pci/cx23885/cx23885-video.c] | | | | | | | | |



Single Instance Design





Automation Design

Automation Design

Major Automated Steps:

- Periodically poll Linux release RSS feed to kick off the process when a new kernel is released
- Download new kernel
- Generate and apply patch
- o Run L-SAP tool
- Run mapping and differencing algorithm on results
- Upload to website

Testing & Evaluation Plan

Testing & Evaluation Plan

- Unit Testing
 - Patching Algorithm, Differencing Algorithm, and The Automation Design are all written in Java
 - Unit Testing will be done through JUnit testing framework
 - Website will be build using TypeScript and the Angular 2 Framework
 - The Karma framework will be used to provide testing to the components on the website
 - Unit Testing will be done every time we intend to merge code with our master branch
- Integration Testing
 - Due to the stages of the pipeline being developed simultaneously, integration testing cannot be done immediately
 - Once the basic components of every component have been created, an Integration Testing project can be built to run through the pipeline (with a mock version of the verifier)

Project Schedule & Management

Project Schedule

- The next semester will be spent implementing existing designs
 - The differencing tool will be completed based on the mapping tool and research from this semester
 - The automation of the patch will be created based on the manual patch created this semester
 - The website will be implemented based on designs from this semester
 - Finally, all the pieces will be tied together for total automation of the tool
 - We will also implement unit testing and integration testing into our development cycle

| Task Week #: | 1 | 2 | 3 | 4 | 5 | 6 | 7 | 8 | 9 | 10 | 11 | 12 | 13 | 14 |
|-------------------|---|---|---|---|---|---|---|---|---|----|----|----|----|----|
| Automate Patch | | | | | | | | | | | | | | |
| Automate Tool | | | | | | | | | | | | | | |
| Diff Algorithm | | | | | | | | | | | | | | |
| Post to Website | | | | | | | | | | | | | | |
| Design Website | | | | | | | | | | | | | | |
| Unit Testing | | | | | | | | | | | | | | |
| Integration Tests | | | | | | | | | | | | | | |

Project Management

- We are using GitLab issues to document work that needs to be done, and to assign it in order to stay on track with our project plan
- Work will be lead in the following ways:
 - Srinivas Dhanwada Differencing algorithm lead
 - Collin McIntyre Patch automation lead
 - Matt Wall Web design lead
 - o Ben Weno Tool automation lead
- Any resources for the project are provided by our client, such as a computer capable of running L-SAP and our tool

Potential Risks:

- A sweeping change in the Linux kernel code could render L-SAP unable to produce proper results
- L-SAP takes a large amount of memory to run, and our tool must run on the same machine, so we will need to keep resource usage low

Summary

Summary

- L-SAP Provides Accurate Results and Human Readable Evidence to aid in verification
- Problems that exist with the current process of running L-SAP
- Our Solution solves those problems and aims to automate the process
 - Patch Creation and Application
 - Difference Mapping and Summarization
 - User-Friendly Website to aid in Understanding of Results
 - o Automation of Patching, Verification and Differencing
- Testing Framework will be integrated into development cycle
- Week-by-week schedule and Project Management

Questions?